

Errata

Coppock, A. (2019). Generalizing from Survey Experiments Conducted on Mechanical Turk: A Replication Approach. *Political Science Research and Methods* 7(3), 613-628.

Five sentences in the published article need correcting. Two state numbers the deposited replication archive does not support, and a third describes the tests' statistical power more generously than the corrected numbers allow. The remaining two are in the footnote that explains the Holm correction: one calls the correction a correlation, and the other states the Benjamini-Hochberg threshold on the false discovery rate in place of the Holm threshold the footnote names and the code applied. All five are in the section reporting the randomization inference tests of treatment effect homogeneity, and none changes a conclusion of the article: the tests still reject the null of constant effects in a small minority of comparisons, the counts in Table 2 come from the code rather than from the footnote's formula, and the correspondence between the Mechanical Turk and probability sample estimates is untouched.

Article DOI: [10.1017/psrm.2018.10](https://doi.org/10.1017/psrm.2018.10). Replication archive: [10.7910/DVN/F1CFFM](https://github.com/active-maintenance-aec/coppock_2019a). The code behind every corrected value below is at https://github.com/active-maintenance-aec/coppock_2019a, and every one of them is recomputed when this document is rendered rather than typed into it.

Corrections are marked in **bold**. Numbers the article hedges with an approximation sign keep the article's own precision, because the approximation is the claim; fuller values appear in the note beneath each entry.

1. Simulated power, page 10

Published. “The MTurk versions of the experiments studied here typically employ 500 subjects per treatment arm, suggesting that we would be well powered (power $\approx$ **0.8**) to detect treatment effect heterogeneity on the scale of 0.2 SD, and moderately powered for 0.1 SD (power $\approx$ **0.6**).”

Corrected. “The MTurk versions of the experiments studied here typically employ 500 subjects per treatment arm, suggesting that we would be **moderately** well powered (power $\approx$ **0.6**) to detect treatment effect heterogeneity on the scale of 0.2 SD, and **poorly** powered for 0.1 SD (power $\approx$ **0.1**).”

The simulation results the article deposits, and which Figure 4 is drawn from, give a rejection probability of 0.61 at 0.2 SD and 0.15 at 0.1 SD, both at 500 subjects per treatment arm. Figure

4 draws guide lines at 500 subjects and at power 0.8 and annotates the 0.2 SD curve where it crosses the second of them, which the simulation puts at 600 subjects per arm rather than 500; the sentence reads the height of the guide line rather than the curve. At those rejection probabilities the test is moderately well powered at 0.2 SD rather than well powered, and poorly powered at 0.1 SD rather than moderately so, and both descriptions are corrected with the numbers. The figure itself is correct and is not reprinted here.

2. Significant heterogeneity tests on the TESS/GfK sample, page 10

Published. “On the TESS/GfK replications, 0 of 10 tests were significant.”

Corrected. “On the TESS/GfK replications, 1 of 10 tests were significant.”

*The deposited output file for Table 2 records 1 significant TESS/GfK test, and so does the maintained pipeline; the count does not move with the change R 3.6.0 made to `sample()`, which moves the other two rows of that column. Table 2 prints 2 in the TESS/GfK row of its *N* Significant column and 1 in the Original row, which is that column read in the row order of the deposited output file rather than in the table’s own; the two entries are interchanged. Correcting the table requires the pre-R-3.6.0 sampler, so it is documented in the repository rather than reprinted here.*

3. The description of statistical power, page 11

Published. “Because we are **relatively** well powered to detect politically meaningful differences in treatment response, I conclude from this evidence that the treatment effect homogeneity explanation of the correspondence across experimental sites is plausible.”

Corrected. “Because we are **moderately** well powered to detect politically meaningful differences in treatment response, I conclude from this evidence that the treatment effect homogeneity explanation of the correspondence across experimental sites is plausible.”

The sentence rests on the power figures of entry 1. At the corrected rejection probabilities, 0.61 at 0.2 SD and 0.15 at 0.1 SD, the tests are moderately well powered rather than relatively well powered, which is the description entry 1 now gives them. The conclusion the sentence draws is unchanged.

4. The word naming the procedure in footnote 8, page 11

Published. “To implement the **correlation**, order the m uncorrected p-values within a “family” from smallest to largest.”

Corrected. “To implement the **correction**, order the m uncorrected p-values within a “family” from smallest to largest.”

The note is about the Holm correction, which its own first sentence names twice, and the sentence that says how to apply it calls it a correlation. Nothing in the analysis turns on the word.

5. The condition stated for the Holm correction, footnote 8, page 11

Published. “Identify the smallest p-value for which the following condition holds: $p_k \leq \frac{k}{m}\alpha$, where k indexes the p-values, and α the target significance level.”

Corrected. “Identify the smallest p-value for which the following condition holds: $p_k > \alpha/(m - k + 1)$, where k indexes the p-values, and α the target significance level.”

The condition the note prints is the Benjamini-Hochberg threshold on the false discovery rate, not the Holm threshold on the family-wise error rate that the note names and that the analysis used. The rest of the note is already Holm’s step-down rule: it says to find the SMALLEST p-value meeting the condition and to treat it and every larger p-value as insignificant, which is where Holm stops and is not how Benjamini-Hochberg is read. So the single defect is the condition, and correcting it makes the note describe the procedure the code ran. Both scripts call `p.adjust(max_pval, "holm")`, the deposit’s own `coppock_generalizability_ri_analysis.R` and the maintained `table_2_ri_heterogeneity.R`, so Table 2’s N Significant (Holm) column and every number under it are unaffected.