

Active Maintenance Report: coppock_2019a

Coppock (2019, PSRM): Generalizing from Survey Experiments Conducted on Mechanical Turk

Table of contents

1	Paper overview	3
2	Summary	3
2.1	Does the deposited archive run?	3
2.2	Does the maintained rewrite reproduce the paper?	4
3	Original archive reproducibility	5
3.1	Filename case	5
3.2	The archive writes into itself	6
3.3	The standard error package no longer exists	6
4	Number-by-number comparison	6
4.1	Table 2 in full	10
5	Maintained rewrite	11
5.1	Deprecated patterns replaced	12
5.2	What the rewrite does not cover	12
5.3	The power sentence	13
5.4	Determinism	13
6	Figures	15
7	Tables	17
8	The extraction and the two instruments	19
9	Maintained rewrite verification	20

Drafted by Claude Opus 5 under the supervision of Alex Coppock.

This repository holds the actively maintained replication code for Coppock (2019), together with the reproducibility report that documents what the original archive did and did not do. It is part of a program applying the maintenance proposal in Peer, Orr and Coppock (2021, *PS: Political Science & Politics*, doi [10.1017/S1049096521000366](https://doi.org/10.1017/S1049096521000366)) to a set of published archives.

Article	10.1017/psrm.2018.10
Replication archive	10.7910/DVN/F1CFFM

The data are not redistributed here. The deposit is 38 files at Harvard Dataverse, which is the only copy this repository points at. `download_original.R` fetches it and verifies every file; `original_manifest.csv` pins the file identifiers, sizes and checksums, so the exact bytes this code was written against are recorded in version control even though the bytes themselves are not.

Repository layout. `maintained/` is the maintained rewrite: one script per published table or figure, writing to `output/`, which is committed so a reader can compare a fresh run against it without downloading anything, plus `in_text_claims.R`, which recomputes every number the article states in prose. `ground_truth/` ties every published number to the code that produces it: `published_claims.csv` is the extraction of the article, `build_ground_truth.R` builds the comparison and gates the pipeline on it. `original/` is created by the download script and is deliberately absent from the repository. This README is the reproducibility report, also available as a PDF in `report/`. `coppock_2019a_errata.pdf` corrects two sentences of the published article.

License. CC0 1.0 Universal, matching the terms of the deposit this repository maintains. See LICENSE.

To reproduce. Clone or download the repository, open `coppock_2019a.Rproj`, and run:

```
source("run_all.R")
```

That fetches the deposit, verifies its 38 files, and produces every table and figure into `maintained/output/`. Required packages: `tidyverse`, `sandwich`, `lmtest`, `estimatr`, `metap`, `knitr`, `kableExtra`, `here`. `metap` depends on `multtest` from Bioconductor, so it needs `BiocManager::install("multtest")` before `install.packages("metap")`. Paths resolve through `here`, so nothing depends on the working directory. A full run takes about a minute, most of it `table_2_ri_heterogeneity.R` recomputing 90 randomization inference tests. A successful run overwrites `maintained/output/`, which is committed: **git diff on that folder is the reproduction check.**

1 Paper overview

Citation: Coppock, A. (2019). “Generalizing from Survey Experiments Conducted on Mechanical Turk: A Replication Approach.” *Political Science Research and Methods*, 7(3), 613-628. DOI: 10.1017/psrm.2018.10

Replication archive: <https://doi.org/10.7910/DVN/F1CFFM>

Summary: The paper asks whether average treatment effects estimated on Mechanical Turk generalize to national probability samples. Fifteen survey experiments, originally fielded on probability or convenience samples, were replicated on MTurk, and seven were replicated a second time on TESS/GfK national probability samples. Treatments are framing, priming and information manipulations; outcomes are attitudes, standardized by the control group standard deviation in the original study. Effects are estimated by difference in means with Neyman (HC2) standard errors. Across 40 treatment-versus-control comparisons the MTurk and original estimates correlate at 0.85, and a randomization inference test of the null of treatment effect homogeneity rarely rejects, which is the mechanism the paper offers for why the samples agree.

2 Summary

Two questions, answered before the detail.

2.1 Does the deposited archive run?

Not as deposited, and the reasons split cleanly into the trivial and the substantive.

The trivial ones are two absent packages and a systematic path error. `metap`, used to combine p-values by Fisher’s method, is not installed by default and pulls `multtest` from Bioconductor. `commarobust`, which supplies every standard error in the paper, was never on CRAN and now cannot be installed in working form at all; that one is treated below, because it has outgrown the trivial category. The path error is that eight `load()` calls name files ending `.rdata` when the deposited files end `.RData`. macOS does not care and Linux does, so six of the archive’s seven scripts fail on their first data load on any case-sensitive filesystem. Nothing about the analysis is wrong; the archive simply cannot start.

The substantive ones are three.

First, two mistyped `sink()` calls silently destroyed two of the archive’s own outputs. At line 249 of `coppock_generalizability_appendix.R`, a `sink()` naming Table 1’s file stands

where a bare `sink()` was meant, so instead of closing the Table 6 sink it reopens Table 1's file and truncates it: the deposited `..._table_1.txt` is zero bytes. At line 566 the same script writes Table 15's results to `..._table_5.txt`, overwriting the Patriot Act table that belongs there. The deposited `..._table_5.txt` contains the terrorism spending results, and no `..._table_15.txt` exists. Both errors are silent, and both survived into the published deposit.

Second, the uncorrected counts in Table 2 do not come back. The cause is not a bug in the archive but the change R 3.6.0 made to `sample()`, which replaced the rounding sampler with the rejection sampler and so changed which permutations a given seed draws. Setting `RNGkind(sample.kind = "Rounding")` and rerunning reproduces the deposited output exactly, 2 and 9 and 1 across the three samples, where the current sampler gives 1 and 8 and 1. The archive is reproducible; it is reproducible under the R of 2018.

Third, and separately from anything R did, the published Table 2's N Significant column is transposed. The deposited output file prints its three rows in alphabetical order of the sample name, TESS/GfK first and original last, and the published table prints them in the reverse order, original first. Table 2's N Comparisons column was reordered to match, 40 and 40 and 10; its N Significant column was not, so the published 1 in the Original row is the TESS/GfK count and the published 2 in the TESS/GfK row is the original count. The correct entries in the published row order are 2, 9 and 1, which is exactly what the deposit produces. The Holm-corrected column, which is the one the argument rests on, is symmetric under the same interchange and reproduces exactly under both samplers.

Fourth, the article's two power figures are not the ones its own simulation gives. The section is treated below and in the errata.

2.2 Does the maintained rewrite reproduce the paper?

For 94 of the 100 comparable claims, yes. The correlations that carry the paper's argument reproduce to the reported precision (0.85, 0.90, 0.96, 0.85), as does the 72.5 percent replication rate. All 90 published Table A1 entries reproduce, each an estimate, a standard error and a significance star; so do all 180 treatment estimate and standard error cells printed across the seventeen online appendix regression tables.

The 6 claims that do not match fall into three groups. Two are the uncorrected significant-test counts for MTurk and TESS/GfK in Table 2, which are sampler-dependent for the reasons above; the rewrite reports what current R produces rather than reproducing the 2018 draw, because pinning a deprecated sampler to recover a specific integer would document the number rather than the finding. Two more are the power figures the article quotes from Figure 4, which its own deposited simulation contradicts. The last two are the sentence stating that no TESS/GfK test was significant, where the deposit says one was, and the sentence saying the analysis is limited to two dependent variables per study, which holds or fails depending on what counts as a study.

The substantive claim, that these tests rarely reject the null of treatment effect homogeneity, holds in every version: at most 9 of 40 comparisons reject before correction and at most 8 after.

3 Original archive reproducibility

Table 1: Original archive reproducibility, checked against R 4.6.0 on a case-sensitive reading of the deposited filenames.

Script	Status on current R	Resolution
coppock_generalizability_ri_functions.R	Clean (sourced by others; defines no paths)	No changes required
coppock_generalizability_analysis.R	Fails on load(); fails again on commarobust	Correct the filename case; inline the retired commarobust call
coppock_generalizability_figures.R	Fails on load()	Correct the filename case
coppock_generalizability_table.R	Fails on load()	Correct the filename case
coppock_generalizability_ri_analysis.R	Fails on load(); fails again on metap and commarobust	Correct the filename case; install multtest and metap
coppock_generalizability_simulation.R	Fails on load()	Correct the filename case
coppock_generalizability_appendix.R	Fails on load(); writes two outputs to the wrong file	Correct the filename case; correct two sink() calls

3.1 Filename case

Every deposited data file ends `.RData`. Every script asks for `.rdata`:

Table 2: The eight load() calls whose case does not match the deposited files. Each deposited counterpart ends `.RData`.

Call site	Filename requested
coppock_generalizability_analysis.R:20	coppock_generalizability_studies.rdata
coppock_generalizability_figures.R:21	coppock_generalizability_analysis_results.rdata
coppock_generalizability_figures.R:22	coppock_generalizability_study_data.rdata
coppock_generalizability_table.R:22	coppock_generalizability_analysis_results.rdata
coppock_generalizability_table.R:23	coppock_generalizability_study_data.rdata
coppock_generalizability_ri_analysis.R:20	coppock_generalizability_studies.rdata
coppock_generalizability_appendix.R:20	coppock_generalizability_studies.rdata
coppock_generalizability_simulation.R:60	coppock_generalizability_simulation_results.rdata

The same defect had propagated into the maintained rewrite, whose six `load()` calls were copied from the archive. They are corrected here.

3.2 The archive writes into itself

Six of the archive’s scripts write their outputs into the archive directory, and five of those outputs are themselves deposited files. Running the deposit in place therefore overwrites part of the deposit: `coppock_generalizability_analysis_results.RData`, the three figure PDFs, and `coppock_generalizability_table_2.txt` all change. A seventh file, `ri_results.rdata`, is created and is not part of the deposit at all. Rerunning `download_original.R` restores everything, which is the reason the manifest records checksums rather than trusting whatever happens to be on disk.

3.3 The standard error package no longer exists

Every estimate in the paper passes through `commarobust::commarobust_tidy()`. `commarobust` was never on CRAN. Its author retired it in November 2018 in favour of `estimatr`; the current GitHub head contains one function, a startup message reading “the commarobust package has been entirely replaced by the estimatr package”, and exports nothing. Installing it from the default branch and calling `commarobust::commarobust` raises “not an exported object”. The last working code sits on a branch named `legacy`.

An archive whose standard errors depend on an unversioned GitHub branch is exposed to a decision nobody in this project controls, so the rewrite does not depend on it. The retired `commarobust(fit)` was exactly

```
lmtest::coeftest(fit, sandwich::vcovHC(fit, type = "HC2"))
```

which is what `estimatr::lm_robust()` computes by default: HC2 is both its default variance estimator and the Neyman variance the paper reports. Every model in `maintained/analysis_ate_estimates.R` is therefore fitted with `lm_robust()` directly, and no helper stands between the call site and the estimator. Reproduction is exact to floating point: across all 90 estimates the largest disagreement with the values the archive’s own package produced is 1.4e-12 in the point estimates and 8.0e-14 in the standard errors.

4 Number-by-number comparison

Table 3: Ground truth: 145 rows, comparing the published value to the deposit’s own committed output.

Location	Quantity	Paper	Script	Match
----------	----------	-------	--------	-------

Text, p. 2	Pairs of estimates entering the correlation	40	40.0000	1
Text, p. 2	MTurk-original correlation	0.85	0.8512	1
Text, p. 8	Original-MTurk pairs of coefficients	40	40.0000	1
Text, p. 8	MTurk-original correlation, restated	0.85	0.8512	1
Text, p. 8	MTurk-original correlation on the TESS/GfK subset	0.90	0.9017	1
Text, p. 8	TESS/GfK-MTurk correlation	0.96	0.9582	1
Text, p. 8	TESS/GfK-original correlation	0.85	0.8517	1
Text, p. 8	Coefficients significant in the original studies	25	25.0000	1
Text, p. 8	Of those, also significant on MTurk	18	18.0000	1
Text, p. 8	Coefficients not significant in the original studies	15	15.0000	1
Text, p. 8	Of those, also not significant on MTurk	11	11.0000	1
Text, p. 8	Replicated significant coefficients, restated inside the formula	18	18.0000	1
Text, p. 8	Replicated null coefficients, restated inside the formula	11	11.0000	1
Text, p. 8	Originally significant coefficients, restated inside the formula	25	25.0000	1
Text, p. 8	Originally null coefficients, restated inside the formula	15	15.0000	1
Text, p. 8	Replication rate as a percentage	72.5	72.5000	1
Abstract	Replication experiments reported	15	15.0000	1
Text, p. 2	Replication studies, restated	15	15.0000	1
Text, p. 10	Replications the heterogeneity test covers	15	15.0000	1
Text, p. 2	Study pairs behind the correlation	12	12.0000	1
Text, p. 8	Study pairs, restated	12	12.0000	1
Results I heading, p. 5	Studies the section covers	12	12.0000	1
Text, p. 10	Original studies in the test	12	12.0000	1
Text, p. 10	Original studies, restated in the sentence	12	12.0000	1
Text, p. 10	Study versions the test covers	27	27.0000	1
Text, p. 8	Coefficients behind the correlation	40	40.0000	1
Text, pp. 5-6	Studies replicated on both platforms	3	3.0000	1
Text, p. 5	Cases with a fresh probability sample	3	3.0000	1
Text, p. 8	Studies with a TESS/GfK version	3	3.0000	1
Text, p. 8	Coefficients in those studies	10	10.0000	1
Text, p. 2	Degrees of freedom of the correlation	38	38.0000	1
Text, p. 8	Degrees of freedom, restated	38	38.0000	1
Text, p. 8	Degrees of freedom on the TESS/GfK subset	8	8.0000	1
Text, p. 8	Degrees of freedom, TESS/GfK against MTurk	8	8.0000	1
Text, p. 8	Degrees of freedom, TESS/GfK against original	8	8.0000	1
Figure 2	Facets, one per study	12	12.0000	1
Figure 2	Estimates plotted	90	90.0000	1
Figure 2	Facets carrying two study versions	9	NA	NA
Figure 2	Facets carrying three study versions	3	NA	NA
Figure 2	Versions in each two-version facet	2	NA	NA
Figure 2	Versions in each three-version facet	3	NA	NA
Figure 2	Confidence level of the plotted intervals	95	NA	NA
Figure 3	Coefficient pairs plotted	40	40.0000	1
Figure 3	Confidence level of the plotted intervals	95	NA	NA
Text, p. 10	Simulation parameters varied	2	NA	NA
Figure 4	Heterogeneity scale named for the high-power claim	0.2	0.2000	1
Figure 4	Heterogeneity scale named for the moderate-power claim	0.1	0.1000	1
Figure 4	Simulated power at 0.2 SD and 500 subjects per arm	0.8	0.6100	0
Figure 4	Simulated power at 0.1 SD and 500 subjects per arm	0.6	0.1500	0
Table 1	Estimates reported for concealed_carry	1	1.0000	1
Table 1	Estimates reported for death_penalty	1	1.0000	1
Table 1	Estimates reported for elite_endorsements	2	2.0000	1
Table 1	Estimates reported for expert_economists	5	5.0000	1
Table 1	Estimates reported for frame_breadth	4	4.0000	1
Table 1	Estimates reported for free_trade	4	4.0000	1
Table 1	Estimates reported for immigration	2	2.0000	1
Table 1	Estimates reported for mental_illness	6	6.0000	1
Table 1	Estimates reported for patriot_act	9	9.0000	1
Table 1	Estimates reported for polarization	2	2.0000	1

Table 1	Estimates reported for superordinate_id	2	2.0000	1
Table 1	Estimates reported for system_threat	2	2.0000	1
Table 1	Subjects in the original Haider-Markel and Joslyn (2001) analysis	518	NA	NA
Table 1	Subjects in the MTurk replication of Haider-Markel and Joslyn (2001)	1009	NA	NA
Table 1	Subjects in the original Brader (2005) analysis	281	NA	NA
Table 1	Subjects in the MTurk replication of Brader (2005)	1709	NA	NA
Table 1	Subjects in the original Peffley and Hurwitz (2007) analysis	905	NA	NA
Table 1	Subjects in the MTurk replication of Peffley and Hurwitz (2007)	1285	NA	NA
Table 1	Subjects in the original Transue (2007) analysis	345	NA	NA
Table 1	Subjects in the MTurk replication of Transue (2007)	367	NA	NA
Table 1	Subjects in the original Chong and Druckman (2010) analysis	1302	NA	NA
Table 1	Subjects in the MTurk replication of Chong and Druckman (2010)	1887	NA	NA
Table 1	Subjects in the original Nicholson (2012) analysis	1491	NA	NA
Table 1	Subjects in the MTurk replication of Nicholson (2012)	1249	NA	NA
Table 1	Subjects in the original McGinty, Webster and Barry (2013) analysis	2935	NA	NA
Table 1	Subjects in the MTurk replication of McGinty, Webster and Barry (2013)	2487	NA	NA
Table 1	Subjects in the original Craig and Richeson (2014) analysis	611	NA	NA
Table 1	Subjects in the MTurk replication of Craig and Richeson (2014)	709	NA	NA
Table 1	Subjects in the original Johnston and Ballard (2016) analysis	2041	NA	NA
Table 1	Subjects in the MTurk replication of Johnston and Ballard (2016)	2985	NA	NA
Table 1	Subjects in the original Hiscox (2006) analysis	1610	NA	NA
Table 1	Subjects in the MTurk replication of Hiscox (2006)	2972	NA	NA
Table 1	Subjects in the TESS/GfK replication of Hiscox (2006)	2084	NA	NA
Table 1	Subjects in the original Levendusky and Malhotra (2016) analysis	1041	NA	NA
Table 1	Subjects in the MTurk replication of Levendusky and Malhotra (2016)	1987	NA	NA
Table 1	Subjects in the TESS/GfK replication of Levendusky and Malhotra (2016)	1411	NA	NA
Table 1	Subjects in the original Hopkins and Mummolo (2017) analysis	3269	NA	NA
Table 1	Subjects in the MTurk replication of Hopkins and Mummolo (2017)	2972	NA	NA
Table 1	Subjects in the TESS/GfK replication of Hopkins and Mummolo (2017)	3189	NA	NA
Table 2	Table 2, original, comparisons tested	40	40.0000	1
Table 2	Table 2, original, significant after the Holm correction	0	0.0000	1
Table 2	Table 2, mt, comparisons tested	40	40.0000	1
Table 2	Table 2, mt, significant after the Holm correction	8	8.0000	1
Table 2	Table 2, gfk, comparisons tested	10	10.0000	1
Table 2	Table 2, gfk, significant after the Holm correction	0	0.0000	1
Table 2	Table 2, original, significant before correction	1	2.0000	0
Table 2	Table 2, MTurk, significant before correction	9	9.0000	1
Table 2	Table 2, TESS/GfK, significant before correction	2	1.0000	0
Text, p. 10	Original comparisons tested, as the text states them	40	40.0000	1
Text, p. 10	MTurk comparisons tested, as the text states them	40	40.0000	1
Text, p. 10	TESS/GfK comparisons tested, as the text states them	10	10.0000	1
Text, p. 10	Original comparisons the text says rejected homogeneity	1	2.0000	0
Text, p. 10	MTurk comparisons the text says rejected homogeneity	8	9.0000	0

Text, p. 10	TESS/GfK comparisons the text says rejected homogeneity	0	1.0000	0
Table A1	Table A1 original entries reproduced of 40 published	40	40.0000	1
Table A1	Table A1 mt entries reproduced of 40 published	40	40.0000	1
Table A1	Table A1 gfk entries reproduced of 10 published	10	10.0000	1
Online Appendix Table 1	Treatment estimate and standard error cells reproduced of 4 published	4	NA	NA
Online Appendix Table 1	Intercept, sample size and R-squared cells of appendix Table 1 (Gun Control)	8	NA	NA
Online Appendix Table 2	Treatment estimate and standard error cells reproduced of 8 published	8	NA	NA
Online Appendix Table 2	Intercept, sample size and R-squared cells of appendix Table 2 (Immigration)	16	NA	NA
Online Appendix Table 3	Treatment estimate and standard error cells reproduced of 4 published	4	NA	NA
Online Appendix Table 3	Intercept, sample size and R-squared cells of appendix Table 3 (Death Penalty)	8	NA	NA
Online Appendix Table 4	Treatment estimate and standard error cells reproduced of 8 published	8	NA	NA
Online Appendix Table 4	Intercept, sample size and R-squared cells of appendix Table 4 (Superordinate Identity)	8	NA	NA
Online Appendix Table 5	Treatment estimate and standard error cells reproduced of 36 published	36	NA	NA
Online Appendix Table 5	Intercept, sample size and R-squared cells of appendix Table 5 (Patriot Act)	8	NA	NA
Online Appendix Table 6	Treatment estimate and standard error cells reproduced of 8 published	8	NA	NA
Online Appendix Table 6	Intercept, sample size and R-squared cells of appendix Table 6 (Elite Endorsements)	24	NA	NA
Online Appendix Table 7	Treatment estimate and standard error cells reproduced of 24 published	24	NA	NA
Online Appendix Table 7	Intercept, sample size and R-squared cells of appendix Table 7 (Mental Illness)	16	NA	NA
Online Appendix Table 8	Treatment estimate and standard error cells reproduced of 8 published	8	NA	NA
Online Appendix Table 8	Intercept, sample size and R-squared cells of appendix Table 8 (System Threat)	16	NA	NA
Online Appendix Table 9	Treatment estimate and standard error cells reproduced of 12 published	12	NA	NA
Online Appendix Table 9	Intercept, sample size and R-squared cells of appendix Table 9 (Expert Economists: Immigration, Health Care and China)	24	NA	NA
Online Appendix Table 10	Treatment estimate and standard error cells reproduced of 8 published	8	NA	NA
Online Appendix Table 10	Intercept, sample size and R-squared cells of appendix Table 10 (Expert Economists: Tax Cuts and Gold Standard)	16	NA	NA
Online Appendix Table 11	Treatment estimate and standard error cells reproduced of 24 published	24	NA	NA
Online Appendix Table 11	Intercept, sample size and R-squared cells of appendix Table 11 (Free Trade)	12	NA	NA
Online Appendix Table 12	Treatment estimate and standard error cells reproduced of 6 published	6	NA	NA
Online Appendix Table 12	Intercept, sample size and R-squared cells of appendix Table 12 (Extremity)	12	NA	NA
Online Appendix Table 13	Treatment estimate and standard error cells reproduced of 6 published	6	NA	NA
Online Appendix Table 13	Intercept, sample size and R-squared cells of appendix Table 13 (Perceived Polarization)	12	NA	NA
Online Appendix Table 14	Treatment estimate and standard error cells reproduced of 6 published	6	NA	NA
Online Appendix Table 14	Intercept, sample size and R-squared cells of appendix Table 14 (Crime Spending)	12	NA	NA

Online Appendix Table 15	Treatment estimate and standard error cells reproduced of 6 published	6	NA	NA
Online Appendix Table 15	Intercept, sample size and R-squared cells of appendix Table 15 (Health Care Spending)	12	NA	NA
Online Appendix Table 16	Treatment estimate and standard error cells reproduced of 6 published	6	NA	NA
Online Appendix Table 16	Intercept, sample size and R-squared cells of appendix Table 16 (Stimulus Spending)	12	NA	NA
Online Appendix Table 17	Treatment estimate and standard error cells reproduced of 6 published	6	NA	NA
Online Appendix Table 17	Intercept, sample size and R-squared cells of appendix Table 17 (Terrorism Spending)	12	NA	NA
Text, p. 6	Most dependent variables in any of the deposit's study groupings	2	5.0000	0
Footnote 5	Mean absolute standardized effect	0.2	0.2002	1
Text, p. 8	Pairs significant in both versions with opposite signs	0	0.0000	1
Text, pp. 10-11	Comparisons rejecting homogeneity in more than one sample	1	NA	NA
Figure 4	Subjects per treatment arm the MTurk replications typically employ	500	NA	NA

67 of the 75 published values with an archive counterpart are reproduced by the deposit's own committed output. `value_script` is read out of the objects and text files the archive itself deposited rather than typed or produced by re-running anything, and it is blank where the deposit prints nothing: Table 1's sample sizes, and the online appendix regression tables, whose deposited text files include two the archive's own mistyped `sink()` calls destroyed.

4.1 Table 2 in full

Table 4: Table 2's counts of significant homogeneity tests, from five sources. 'Deposited output' is the archive's own table 2 text file as it sits in the deposit. 'Rounding sampler' restores the pre-R-3.6.0 behaviour of the sampler behind `sample()`.

Cell	Table 2	Body text	Deposited output	Rounding sampler, 2026	Current sampler, 2026
Original, uncorrected	1	1	2	2	1
Mechanical Turk, uncorrected	9	8	9	9	8
TESS/GfK, uncorrected	2	0	1	1	1
Original, Holm	0		0	0	0
Mechanical Turk, Holm	8		8	8	8
TESS/GfK, Holm	0	0	0	0	0

Three things follow. The archive is internally reproducible: the rounding sampler recovers the deposited output in every cell. The R 3.6.0 sampler change moves two cells, the MTurk and original counts, because one test in each of those samples has a maximum p-value of exactly 0.05 under the 2018 draw and none does under the current one.

And the published table's N Significant column is the deposited column in the deposited file's own row order. That file prints its rows alphabetically, `gfk` then `mt` then `original`, and reads 1, 9, 2 down the column. The published table prints its rows in the reverse order, Original

then Mechanical Turk then TESS/GfK, and its N Significant column also reads 1, 9, 2. The N Comparisons column beside it was reordered correctly, 40 and 40 and 10 rather than 10 and 40 and 40, which is why the interchange is invisible from inside the table. The Holm column is symmetric under it and so cannot show it either. Under this reading every cell of the published table is accounted for: the correct entries in the published row order are 2, 9 and 1.

The alternative, that Table 2 was typeset from a run other than the one deposited, does not survive measurement. `ground_truth/measure_table_2_dispersion.R` runs the deposited procedure over a range of seeds under both samplers and records what it gets:

Table 5: Counts of significant homogeneity tests across 24 runs of the deposited procedure at different seeds.

Sample	Sampler	Runs	N Significant, values taken
gfk	rejection	16	1
gfk	rounding	8	1
mt	rejection	16	8, 9, 10
mt	rounding	8	8, 9, 10
original	rejection	16	1, 2
original	rounding	8	1, 2

The TESS/GfK count was 1 in every one of the 24 runs and never 2, while the original and MTurk counts both moved. A published TESS/GfK count of 2 is not something the deposited code produces.

The body text is a separate matter and is not settled by any of this. It reads 1, 8 and 0 where the deposit reads 2, 9 and 1. Its first two figures are the counts a strict < 0.05 would give on the deposit's own draw, where Table 2's own rule is $\leq$, so the sentence and the table can differ legitimately by one in each of those samples. Its third figure has no such account: the TESS/GfK count is 1 under either rule and under either sampler. That sentence is corrected in `coppock_2019a_errata.pdf`; the other two are recorded as unresolved.

5 Maintained rewrite

`maintained/` holds eight scripts and a shared `helpers.R`, replacing the archive's six analysis scripts and its function file. The rewrite is a translation: estimators, specifications and sample restrictions are unchanged.

Table 6: Maintained rewrite script inventory.

Script	Output
helpers.R	packages; Table A1 cell formatter
analysis_ate_estimates.R	output/analysis_results.rds
table_a1_ate_summary.R	output/table_a1_ate_summary.csv, .tex
table_2_ri_heterogeneity.R	output/table_2_ri_heterogeneity.csv, .tex, output/ri_results.rds
figure_2_study_estimates.R	output/figure_2_study_estimates.csv, .pdf, .png
figure_3_mt_original_scatter.R	output/figure_3_mt_original_scatter.csv, .pdf, .png
figure_4_power_simulation.R	output/figure_4_power_simulation.csv, .pdf, .png
text_correlations.R	output/text_correlations.csv
in_text_claims.R	printed CLAIM lines only

5.1 Deprecated patterns replaced

Table 7: Deprecated patterns and their replacements in the maintained rewrite.

Original pattern	Replacement
<code>rm(list = ls())</code>	(omitted)
<code>setwd("\")</code>	<code>here::here()</code>
<code>load("\...rdata\")</code> against a ‘RData’ file	the deposited filename, exact case
<code>commarobust::commarobust_tidy()</code>	<code>estimatr::lm_robust()</code> , whose default is HC2
<code>do(...\$summary_df)</code>	<code>reframe(...)</code> with <code>pick(everything())</code>
<code>spread()</code> / <code>gather()</code>	<code>pivot_wider()</code> / <code>pivot_longer()</code>
<code>geom_errorbarh()</code>	<code>geom_linerange()</code>
<code>ggstance::position_dodgev()</code>	<code>position_dodge(width = 0.5)</code>
<code>reshape2::melt()</code>	(the archive’s saved result is read directly)
<code>sink()</code> to a hardcoded path	<code>write_csv()</code> and <code>readr::write_lines()</code> to ‘output/’
magrittr pipes	the native pipe

5.2 What the rewrite does not cover

The archive’s `coppock_generalizability_appendix.R` produces the seventeen per-study regression tables of the online appendix. Each prints the same treatment coefficients the rewrite estimates, to three decimal places rather than Table A1’s two, and every one of those 180 cells reproduces. Each table also prints an intercept, a sample size and an R-squared, and appendix Table 6 prints a party identification covariate. The rewrite drops the intercept and that covariate before writing its estimates and exports neither sample sizes nor R-squared, so those cells have no counterpart in `maintained/output/`; the ground truth records them as uncovered rather than as reproduced. Table 1’s sample size columns are uncovered for the same reason.

Figure 1 is uncovered by anything. It illustrates the theoretical argument with a simulated population, and no script in the deposit draws it.

The rewrite also does not re-run the power simulation behind Figure 4. Neither does the archive: the simulation is commented out in `coppock_generalizability_simulation.R`, with a note that it “takes a *very* long time to run”, and the deposited result object is loaded instead. The figure is reproduced from that object, and `figure_4_power_simulation.R` now writes the plotted grid to a CSV so the numbers the text quotes from it have a source.

5.3 The power sentence

The article says of Figure 4 that at 500 subjects per treatment arm “we would be well powered (power ≈ 0.8) to detect treatment effect heterogeneity on the scale of 0.2 SD, and moderately powered for 0.1 SD (power ≈ 0.6).”

Table 8: The two power figures the article quotes, against the simulation results it deposits, both at 500 subjects per treatment arm.

Heterogeneity scale	Power stated in the text	Power in the deposited simulation
0.2 SD	0.8	0.61
0.1 SD	0.6	0.15

Figure 4 draws a dashed guide line at 500 subjects and another at power 0.8, and annotates the 0.2 SD curve where it crosses the second of them. The curve first reaches 0.8 at 600 subjects per arm, not 500, so the sentence reads the height of the guide line rather than the curve. The figure is right; the sentence is not, and it is corrected in `coppock_2019a_errata.pdf`. The claim the paper rests on, that the tests are powered against politically meaningful heterogeneity, survives at 0.2 SD and is weaker at 0.1 SD than the text says.

5.4 Determinism

`table_2_ri_heterogeneity.R` sets the archive’s seed once and recomputes all 90 tests on every run rather than reading back a cached result. Two independent runs, five months apart on different R versions, agree in every one of the 90 p-values to the last bit. The 90 tests take a minute. The script deliberately does not skip its work when `output/ri_results.rds` already exists: a script that does that cannot be checked by diffing its own output, which is what this repository asks a reader to do.

6 Figures




Figure 1: Figure 2: standardized ATE estimates by study and sample.




Figure 2: Figure 3: MTurk versus original standardized estimates.




Figure 3: Figure 4: power of the homogeneity test by sample size and effect heterogeneity.

7 Tables

Table 9: Table 2 as the maintained rewrite produces it.

Sample	N comparisons	N significant	N significant (Holm)
gfk	10	1	0
mt	40	8	8
original	40	1	0

Table 10: Table A1 as the maintained rewrite produces it: standardized ATE estimates by study, outcome and sample, with HC2 standard errors in parentheses.

Study	Outcome	Coefficient	Original	MTurk	GfK
Haider-Markel and Joslyn (2001)	Support for Concealed Carry Law	Citizens' Rights Frame	0.33 (0.08)*	0.21 (0.05)*	NA
Brader (2005)	Negative Impact	Positive Frame	-0.28 (0.13)*	-0.22 (0.04)*	NA
NA	Support for Immigration	Positive Frame	0.39 (0.14)*	0.26 (0.05)*	NA
Peffley and Hurwitz (2007)	Favor Death Penalty	African Americans	0.42 (0.09)*	0.10 (0.05)*	NA
Transue (2007)	Willingness to Pay Tax	Other Americans	0.18 (0.10)	-0.08 (0.10)	NA
NA	Willingness to Pay Tax	Public Schools	0.16 (0.10)	0.19 (0.09)*	NA
Chong and Druckman (2010)	Patriot Act Support	Con / Memory Based	-0.38 (0.12)*	-0.18 (0.11)	NA
NA	Patriot Act Support	Con / No Processing	-0.41 (0.13)*	-0.39 (0.11)*	NA
NA	Patriot Act Support	Con / Online Processing	-0.45 (0.12)*	-0.40 (0.11)*	NA
NA	Patriot Act Support	Pro / Memory Based	0.33 (0.12)*	0.20 (0.11)	NA
NA	Patriot Act Support	Pro / No Processing	0.43 (0.12)*	0.32 (0.11)*	NA
NA	Patriot Act Support	Pro / Online Processing	0.33 (0.12)*	0.27 (0.11)*	NA
NA	Patriot Act Support	Both / Memory Based	-0.07 (0.14)	-0.10 (0.12)	NA
NA	Patriot Act Support	Both / No Processing	-0.08 (0.16)	0.01 (0.13)	NA
NA	Patriot Act Support	Both / Online Processing	0.01 (0.13)	-0.16 (0.13)	NA
Nicholson (2012)	Support for Foreclosure Bill	In Party Cue	0.18 (0.08)*	0.11 (0.07)	NA
NA	Support for Immigration Bill	In Party Cue	0.23 (0.09)*	0.10 (0.06)	NA
McGinty, Webster, and Barry (2013)	Magazines	News	0.11 (0.04)*	0.01 (0.04)	NA
NA	Magazines	LCM Ban	0.18 (0.06)*	0.12 (0.05)*	NA
NA	Magazines	Mental Illness	0.10 (0.05)	-0.01 (0.05)	NA
NA	SMI Danger	News	0.09 (0.05)	0.02 (0.04)	NA
NA	SMI Danger	LCM Ban	-0.07 (0.06)	0.01 (0.05)	NA
NA	SMI Danger	Mental Illness	-0.12 (0.06)*	-0.08 (0.05)	NA
Craig and Richeson (2014)	Support for Immigration	Majority Minority	-0.00 (0.10)	-0.22 (0.07)*	NA
NA	Way of Life	Majority Minority	-0.14 (0.11)	0.08 (0.07)	NA
Johnston and Ballard (2016)	Agree on Gold Standard	Expert Treatment	0.38 (0.11)*	0.50 (0.04)*	NA
NA	Agree on Health Care	Expert Treatment	0.15 (0.10)	0.34 (0.04)*	NA
NA	Agree on Immigration	Expert Treatment	0.14 (0.11)	0.33 (0.04)*	NA
NA	Agree on Tax Cut	Expert Treatment	0.35 (0.10)*	0.44 (0.04)*	NA
NA	Agree on Trade with China	Expert Treatment	0.24 (0.10)*	0.42 (0.04)*	NA
Hiscox (2006)	Support for Free Trade	Expert	0.25 (0.06)*	0.28 (0.03)*	0.17 (0.05)*
NA	Support for Free Trade	Positive	-0.17 (0.08)*	-0.08 (0.04)	-0.03 (0.07)
NA	Support for Free Trade	Negative	-0.28 (0.08)*	-0.32 (0.05)*	-0.32 (0.07)*
NA	Support for Free Trade	Pos + Neg	-0.41 (0.09)*	-0.33 (0.05)*	-0.23 (0.07)*
Levendusky and Malhotra (2015)	Extremity of Policy Views	Polarized Treatment	0.02 (0.08)	-0.05 (0.04)	0.05 (0.06)
NA	Perceived Polarization	Polarized Treatment	0.16 (0.08)*	0.36 (0.05)*	0.47 (0.08)*
Hopkins and Mummolo (2017)	Support for Crime Spending	Crime Argument	0.06 (0.05)	-0.01 (0.05)	0.09 (0.05)
NA	Support for Health Care Spending	Health Care Argument	-0.03 (0.05)	-0.05 (0.04)	-0.03 (0.05)
NA	Support for Stimulus Spending	Stimulus Argument	-0.16 (0.04)*	-0.21 (0.04)*	-0.23 (0.04)*
NA	Support for Terrorism Spending	Terrorism Argument	0.25 (0.05)*	0.11 (0.05)*	0.18 (0.06)*

8 The extraction and the two instruments

Every published number in this article and its online appendix is written down once, in `ground_truth/published_claims.csv`. That file is the extraction: 204 rows, each carrying a claim identifier, where the number appears, what kind of claim it is, the string the page prints and the number of decimals it prints it to. It is the only place in the repository a published number is typed, and nothing in `maintained/` reads it except for those decimal counts.

Table 11: The extraction: 204 published claims, classified by hand.

Claim type	Needs a claim block	Claims
definitional	no	35
definitional	yes	9
descriptive	no	1
descriptive	yes	4
pipeline	no	44
pipeline	yes	65
structural	no	16
structural	yes	22
transcribed	no	8

A **pipeline** claim is one the analysis produces; **descriptive** claims are about shape, sign or count and have no printed number to compare; **definitional** and **structural** claims are design parameters and layout facts; **transcribed** claims are copied from other articles and cannot drift. The **needs_block** column is TRUE where the pipeline can print something true about the claim, which is 100 of the 204. The remainder are verified where they are used: the confidence level of the plausible-effects interval is visible in the `qnorm(0.9999)` call the archive’s own function file makes, Figure 4’s guide lines in the `geom_hline()` and `geom_vline()` that draw them, and the significance threshold behind Table A1’s stars in `helpers.R`.

Two separate instruments then check the same claims.

`ground_truth/build_ground_truth.R` assembles the comparison table. It reads `value_rewrite` back out of `maintained/output/` and `value_script` out of the objects and text files the archive deposited, so neither column can drift from what it describes, and it compares each against the published string by printing the computed number to the page’s own precision rather than by a numeric tolerance, because a double does not record how many decimals the article printed.

`maintained/in_text_claims.R` carries one block per claim needing one: the article’s sentence verbatim in a comment, then code that recomputes the number from the same committed outputs by its own path. It never reads the ground truth. The two reach the same quantities differently, so a disagreement between them is a finding rather than a formality: the ground

truth reads the correlations out of `text_correlations.csv` while the claims file rebuilds them from the estimates.

The build script is the gate. It runs `in_text_claims.R` non-interactively into its own environment, counts the `CLAIM` lines it printed, and stops unless that count equals the 100 claims the extraction says need one, every identifier appears on both sides, every `pipeline` and `descriptive` claim has a ground truth row, every published float has coverage or a stated reason, every adverse verdict carries a `defect_locus` and no clean match does, and the two instruments agree value by value. A block that errored, or that ended in a bare expression and so printed nothing, would satisfy a scan for markers while checking nothing, which is why the file is run rather than read.

Two sentences of the published article are corrected in `coppock_2019a_errata.pdf` at the root of this repository. Neither changes a conclusion.

9 Maintained rewrite verification

Table 12: Maintained rewrite verification: published value against rewrite output.

Location	Quantity	Paper	Rewrite	Match
Text, p. 2	Pairs of estimates entering the correlation	40	40.0000	1
Text, p. 2	MTurk-original correlation	0.85	0.8512	1
Text, p. 8	Original-MTurk pairs of coefficients	40	40.0000	1
Text, p. 8	MTurk-original correlation, restated	0.85	0.8512	1
Text, p. 8	MTurk-original correlation on the TESS/GfK subset	0.90	0.9017	1
Text, p. 8	TESS/GfK-MTurk correlation	0.96	0.9582	1
Text, p. 8	TESS/GfK-original correlation	0.85	0.8517	1
Text, p. 8	Coefficients significant in the original studies	25	25.0000	1
Text, p. 8	Of those, also significant on MTurk	18	18.0000	1
Text, p. 8	Coefficients not significant in the original studies	15	15.0000	1
Text, p. 8	Of those, also not significant on MTurk	11	11.0000	1
Text, p. 8	Replicated significant coefficients, restated inside the formula	18	18.0000	1
Text, p. 8	Replicated null coefficients, restated inside the formula	11	11.0000	1
Text, p. 8	Originally significant coefficients, restated inside the formula	25	25.0000	1
Text, p. 8	Originally null coefficients, restated inside the formula	15	15.0000	1
Text, p. 8	Replication rate as a percentage	72.5	72.5000	1
Abstract	Replication experiments reported	15	15.0000	1
Text, p. 2	Replication studies, restated	15	15.0000	1
Text, p. 10	Replications the heterogeneity test covers	15	15.0000	1
Text, p. 2	Study pairs behind the correlation	12	12.0000	1
Text, p. 8	Study pairs, restated	12	12.0000	1
Results I heading, p. 5	Studies the section covers	12	12.0000	1
Text, p. 10	Original studies in the test	12	12.0000	1
Text, p. 10	Original studies, restated in the sentence	12	12.0000	1
Text, p. 10	Study versions the test covers	27	27.0000	1
Text, p. 8	Coefficients behind the correlation	40	40.0000	1
Text, pp. 5-6	Studies replicated on both platforms	3	3.0000	1
Text, p. 5	Cases with a fresh probability sample	3	3.0000	1
Text, p. 8	Studies with a TESS/GfK version	3	3.0000	1
Text, p. 8	Coefficients in those studies	10	10.0000	1

Text, p. 2	Degrees of freedom of the correlation	38	38.0000	1
Text, p. 8	Degrees of freedom, restated	38	38.0000	1
Text, p. 8	Degrees of freedom on the TESS/GfK subset	8	8.0000	1
Text, p. 8	Degrees of freedom, TESS/GfK against MTurk	8	8.0000	1
Text, p. 8	Degrees of freedom, TESS/GfK against original	8	8.0000	1
Figure 2	Facets, one per study	12	12.0000	1
Figure 2	Estimates plotted	90	90.0000	1
Figure 2	Facets carrying two study versions	9	9.0000	1
Figure 2	Facets carrying three study versions	3	3.0000	1
Figure 2	Versions in each two-version facet	2	2.0000	1
Figure 2	Versions in each three-version facet	3	3.0000	1
Figure 2	Confidence level of the plotted intervals	95	95.0004	1
Figure 3	Coefficient pairs plotted	40	40.0000	1
Figure 3	Confidence level of the plotted intervals	95	95.0004	1
Text, p. 10	Simulation parameters varied	2	2.0000	1
Figure 4	Heterogeneity scale named for the high-power claim	0.2	0.2000	1
Figure 4	Heterogeneity scale named for the moderate-power claim	0.1	0.1000	1
Figure 4	Simulated power at 0.2 SD and 500 subjects per arm	0.8	0.6100	0
Figure 4	Simulated power at 0.1 SD and 500 subjects per arm	0.6	0.1500	0
Table 1	Estimates reported for concealed_carry	1	1.0000	1
Table 1	Estimates reported for death_penalty	1	1.0000	1
Table 1	Estimates reported for elite_endorsements	2	2.0000	1
Table 1	Estimates reported for expert_economists	5	5.0000	1
Table 1	Estimates reported for frame_breadth	4	4.0000	1
Table 1	Estimates reported for free_trade	4	4.0000	1
Table 1	Estimates reported for immigration	2	2.0000	1
Table 1	Estimates reported for mental_illness	6	6.0000	1
Table 1	Estimates reported for patriot_act	9	9.0000	1
Table 1	Estimates reported for polarization	2	2.0000	1
Table 1	Estimates reported for superordinate_id	2	2.0000	1
Table 1	Estimates reported for system_threat	2	2.0000	1
Table 1	Subjects in the original Haider-Markel and Joslyn (2001) analysis	518	NA	NA
Table 1	Subjects in the MTurk replication of Haider-Markel and Joslyn (2001)	1009	NA	NA
Table 1	Subjects in the original Brader (2005) analysis	281	NA	NA
Table 1	Subjects in the MTurk replication of Brader (2005)	1709	NA	NA
Table 1	Subjects in the original Peffley and Hurwitz (2007) analysis	905	NA	NA
Table 1	Subjects in the MTurk replication of Peffley and Hurwitz (2007)	1285	NA	NA
Table 1	Subjects in the original Transue (2007) analysis	345	NA	NA
Table 1	Subjects in the MTurk replication of Transue (2007)	367	NA	NA
Table 1	Subjects in the original Chong and Druckman (2010) analysis	1302	NA	NA
Table 1	Subjects in the MTurk replication of Chong and Druckman (2010)	1887	NA	NA
Table 1	Subjects in the original Nicholson (2012) analysis	1491	NA	NA
Table 1	Subjects in the MTurk replication of Nicholson (2012)	1249	NA	NA
Table 1	Subjects in the original McGinty, Webster and Barry (2013) analysis	2935	NA	NA
Table 1	Subjects in the MTurk replication of McGinty, Webster and Barry (2013)	2487	NA	NA
Table 1	Subjects in the original Craig and Richeson (2014) analysis	611	NA	NA
Table 1	Subjects in the MTurk replication of Craig and Richeson (2014)	709	NA	NA
Table 1	Subjects in the original Johnston and Ballard (2016) analysis	2041	NA	NA
Table 1	Subjects in the MTurk replication of Johnston and Ballard (2016)	2985	NA	NA
Table 1	Subjects in the original Hiscox (2006) analysis	1610	NA	NA

Table 1	Subjects in the MTurk replication of Hiscox (2006)	2972	NA	NA
Table 1	Subjects in the TESS/GfK replication of Hiscox (2006)	2084	NA	NA
Table 1	Subjects in the original Levendusky and Malhotra (2016) analysis	1041	NA	NA
Table 1	Subjects in the MTurk replication of Levendusky and Malhotra (2016)	1987	NA	NA
Table 1	Subjects in the TESS/GfK replication of Levendusky and Malhotra (2016)	1411	NA	NA
Table 1	Subjects in the original Hopkins and Mummolo (2017) analysis	3269	NA	NA
Table 1	Subjects in the MTurk replication of Hopkins and Mummolo (2017)	2972	NA	NA
Table 1	Subjects in the TESS/GfK replication of Hopkins and Mummolo (2017)	3189	NA	NA
Table 2	Table 2, original, comparisons tested	40	40.0000	1
Table 2	Table 2, original, significant after the Holm correction	0	0.0000	1
Table 2	Table 2, mt, comparisons tested	40	40.0000	1
Table 2	Table 2, mt, significant after the Holm correction	8	8.0000	1
Table 2	Table 2, gfk, comparisons tested	10	10.0000	1
Table 2	Table 2, gfk, significant after the Holm correction	0	0.0000	1
Table 2	Table 2, original, significant before correction	1	1.0000	1
Table 2	Table 2, MTurk, significant before correction	9	8.0000	0
Table 2	Table 2, TESS/GfK, significant before correction	2	1.0000	0
Text, p. 10	Original comparisons tested, as the text states them	40	40.0000	1
Text, p. 10	MTurk comparisons tested, as the text states them	40	40.0000	1
Text, p. 10	TESS/GfK comparisons tested, as the text states them	10	10.0000	1
Text, p. 10	Original comparisons the text says rejected homogeneity	1	1.0000	1
Text, p. 10	MTurk comparisons the text says rejected homogeneity	8	8.0000	1
Text, p. 10	TESS/GfK comparisons the text says rejected homogeneity	0	1.0000	0
Table A1	Table A1 original entries reproduced of 40 published	40	40.0000	1
Table A1	Table A1 mt entries reproduced of 40 published	40	40.0000	1
Table A1	Table A1 gfk entries reproduced of 10 published	10	10.0000	1
Online Appendix Table 1	Treatment estimate and standard error cells reproduced of 4 published	4	4.0000	1
Online Appendix Table 1	Intercept, sample size and R-squared cells of appendix Table 1 (Gun Control)	8	NA	NA
Online Appendix Table 2	Treatment estimate and standard error cells reproduced of 8 published	8	8.0000	1
Online Appendix Table 2	Intercept, sample size and R-squared cells of appendix Table 2 (Immigration)	16	NA	NA
Online Appendix Table 3	Treatment estimate and standard error cells reproduced of 4 published	4	4.0000	1
Online Appendix Table 3	Intercept, sample size and R-squared cells of appendix Table 3 (Death Penalty)	8	NA	NA
Online Appendix Table 4	Treatment estimate and standard error cells reproduced of 8 published	8	8.0000	1
Online Appendix Table 4	Intercept, sample size and R-squared cells of appendix Table 4 (Superordinate Identity)	8	NA	NA
Online Appendix Table 5	Treatment estimate and standard error cells reproduced of 36 published	36	36.0000	1
Online Appendix Table 5	Intercept, sample size and R-squared cells of appendix Table 5 (Patriot Act)	8	NA	NA
Online Appendix Table 6	Treatment estimate and standard error cells reproduced of 8 published	8	8.0000	1
Online Appendix Table 6	Intercept, sample size and R-squared cells of appendix Table 6 (Elite Endorsements)	24	NA	NA
Online Appendix Table 7	Treatment estimate and standard error cells reproduced of 24 published	24	24.0000	1
Online Appendix Table 7	Intercept, sample size and R-squared cells of appendix Table 7 (Mental Illness)	16	NA	NA
Online Appendix Table 8	Treatment estimate and standard error cells reproduced of 8 published	8	8.0000	1

Online Appendix Table 8	Intercept, sample size and R-squared cells of appendix Table 8 (System Threat)	16	NA	NA
Online Appendix Table 9	Treatment estimate and standard error cells reproduced of 12 published	12	12.0000	1
Online Appendix Table 9	Intercept, sample size and R-squared cells of appendix Table 9 (Expert Economists: Immigration, Health Care and China)	24	NA	NA
Online Appendix Table 10	Treatment estimate and standard error cells reproduced of 8 published	8	8.0000	1
Online Appendix Table 10	Intercept, sample size and R-squared cells of appendix Table 10 (Expert Economists: Tax Cuts and Gold Standard)	16	NA	NA
Online Appendix Table 11	Treatment estimate and standard error cells reproduced of 24 published	24	24.0000	1
Online Appendix Table 11	Intercept, sample size and R-squared cells of appendix Table 11 (Free Trade)	12	NA	NA
Online Appendix Table 12	Treatment estimate and standard error cells reproduced of 6 published	6	6.0000	1
Online Appendix Table 12	Intercept, sample size and R-squared cells of appendix Table 12 (Extremity)	12	NA	NA
Online Appendix Table 13	Treatment estimate and standard error cells reproduced of 6 published	6	6.0000	1
Online Appendix Table 13	Intercept, sample size and R-squared cells of appendix Table 13 (Perceived Polarization)	12	NA	NA
Online Appendix Table 14	Treatment estimate and standard error cells reproduced of 6 published	6	6.0000	1
Online Appendix Table 14	Intercept, sample size and R-squared cells of appendix Table 14 (Crime Spending)	12	NA	NA
Online Appendix Table 15	Treatment estimate and standard error cells reproduced of 6 published	6	6.0000	1
Online Appendix Table 15	Intercept, sample size and R-squared cells of appendix Table 15 (Health Care Spending)	12	NA	NA
Online Appendix Table 16	Treatment estimate and standard error cells reproduced of 6 published	6	6.0000	1
Online Appendix Table 16	Intercept, sample size and R-squared cells of appendix Table 16 (Stimulus Spending)	12	NA	NA
Online Appendix Table 17	Treatment estimate and standard error cells reproduced of 6 published	6	6.0000	1
Online Appendix Table 17	Intercept, sample size and R-squared cells of appendix Table 17 (Terrorism Spending)	12	NA	NA
Text, p. 6	Most dependent variables in any of the deposit's study groupings	2	5.0000	0
Footnote 5	Mean absolute standardized effect	0.2	0.2002	1
Text, p. 8	Pairs significant in both versions with opposite signs	0	0.0000	1
Text, pp. 10-11	Comparisons rejecting homogeneity in more than one sample	1	1.0000	1
Figure 4	Subjects per treatment arm the MTurk replications typically employ	500	NA	NA

94 of the **100** comparable claims match the published values. The 6 that do not are discussed above: Simulated power at 0.2 SD and 500 subjects per arm; Simulated power at 0.1 SD and 500 subjects per arm; Table 2, MTurk, significant before correction; Table 2, TESS/GfK, significant before correction; TESS/GfK comparisons the text says rejected homogeneity; Most dependent variables in any of the deposit's study groupings.

The rows with no `match_rewrite` are those the rewrite does not compute, and they are recorded rather than dropped: Table 1's sample sizes, the intercept, sample size and R-squared rows of the seventeen online appendix tables, and the article's claim about the typical number of subjects per treatment arm.

10 R environment

Item	Value
R version	4.6.0
Platform	aarch64-apple-darwin23
Date run	2026-08-10
tidyverse	2.0.0
sandwich	3.1.1
lmtest	0.9.40
estimatr	1.0.6
metap	1.14
knitr	1.51
kableExtra	1.4.0
here	1.0.2